

AGRICULTURE LEAF DISEASE CLASSIFICATION USING DEEP LEARNING

A PROJECT REPORT

Submitted by

RAVEEN D (960121243038)

DEEPAK M (960121243012)

RACHEL D (960121243035)

ANUSHA S (960121243006)

In Partial fulfilment for the award of the degree of

BACHELOR OF TECHNOLOGY

IN

Artificial Intelligence and Data Science

ANNAI VAILANKANNI COLLEGE OF ENGINEERING

KANYAKUMARI



ANNA UNIVERSITY::CHENNAI 600 025

MAY 2025

BONAFIDE CERTIFICATE

Certified that this project report “**AGRICULTURE LEAF DISEASE CLASSIFICATION USING DEEP LEARNING**” is the bonafide work of **RAVEEND (960121243038), DEEPAK M (960121243012), RACHEL D (960121243035), ANUSHA S (960121243006)** who carried out the project work under my supervision.

SIGNATURE

Mrs.B.JENEFA M.TECH,

HOD/DEPT OF AI&DS

Professor and Head

Department of AI&DS,

Annai Vailankanni College

of Engineering,

Pottalkulam,kanyakumari.

SIGNATURE

Mrs.S.SANGEETHA

M.TECH,

Supervisor,

Assistant Professor,

Department of AI&DS,

Annai Vailankanni college

of Engineering,

Pottalkulam,kanyakumari.

Submitted For the Project Viva-Voce held on _____
at **Annai Vailankanni College of Engineering, pottalkulam, kanyakumari.**

INTERNAL EXAMINER

EXTERNAL EXAMINER

ACKNOWLEDGEMENT

I am deeply conscious of and grateful for the support and guidance of our Honourable Chairman, our respected Dr. P. Peter Jesudhas, my Principal Dr. R. Angeline Prabhavathy, and our Vice Principal Dr. N. Abilash, who have continuously inspired me to reach greater heights in the pursuit of knowledge.

We proudly express our heartfelt gratitude to our esteemed Head of the Department, Mrs. B. Jenefa, for her constant encouragement and valuable assistance throughout the successful completion of this project.

Special thanks must be extended to our internal guide, Mrs. B. JENEFA, M.TECH, Assistant Professor, Department of Artificial Intelligence and Data Science, for her expert advice, valuable insights, and dedicated guidance throughout the entire duration of the project.

We also sincerely thank our project coordinator, guide, and all the faculty members of the Artificial Intelligence and Data Science Department who supported and assisted us in the successful completion of this project.

TABLE OF CONTENTS

S. No.	Chapter Title	Page No.
1	Introduction	1
2	Problem Statement and Objectives	4
3	Literature Survey and Existing System	7
4	Proposed System and Architecture	11
5	Dataset and Image Preprocessing	16
6	Model Selection and Architecture	21
7	Model Training, Metrics, Results and Comparison	26
	7.1 Model Training	
	7.2 Evaluation Metrics	
	7.3 Result Analysis	
	7.4 Comparison with Existing Models	
8	System Advantages, Limitations and Future Scope	30
	8.1 System Advantages	
	8.2 Limitations	
	8.3 Future Scope	
9	References	36
10	Conclusion and Appendix	37
	10.1 Conclusion	
	10.2 Appendix (Code, Graphs, Confusion Matrix, Screenshots)	

ABSTRACT

Agriculture is a critical sector for the economy of many countries, directly influencing food security and farmer livelihoods. However, crop diseases pose a major threat to agricultural productivity by reducing yields and causing significant financial losses. Early and accurate identification of plant diseases is vital to mitigate these adverse effects and to support timely intervention. This project focuses on developing an automated system for classifying crop leaf diseases using deep learning techniques, thereby aiding farmers with precise and rapid disease diagnosis. The deep learning model developed in this project is based on the MobileNetV2 architecture, renowned for its lightweight structure and computational efficiency. MobileNetV2 is particularly suitable for real-time applications and resource-constrained environments such as rural agricultural settings where access to powerful hardware may be limited. The dataset utilized consists of 68,187 high-quality images, representing various classes of healthy and diseased crop leaves. These images were split into training (28,330), testing (28,330), and validation (11,527) subsets to ensure robust model development and evaluation. The model was trained to classify leaf images into their respective disease categories with a high degree of accuracy. After rigorous training and validation, the MobileNetV2-based model achieved an overall accuracy of 92.45%. This performance demonstrates its reliability in detecting multiple crop diseases from leaf images, making it a practical tool for agricultural disease management. To maximize accessibility, the trained model was integrated into a user-friendly web application. This platform allows users, particularly farmers, to upload leaf images and receive immediate disease classification results. Such instant feedback facilitates early disease detection and enables timely treatment decisions, potentially preventing extensive crop damage and financial loss. The project highlights the potential of combining artificial intelligence and agriculture to create innovative, cost-effective solutions for crop health monitoring. By leveraging the efficiency of MobileNetV2, this system is suitable for deployment on mobile and embedded devices, making it scalable and accessible for widespread agricultural use. Ultimately, this technology contributes to improved crop health, enhanced productivity, and sustainable agricultural practices.

CHAPTER 1

INTRODUCTION

Agriculture plays a vital role in the global economy, especially in developing countries like India, where it accounts for about 18% of the Gross Domestic Product (GDP) and provides employment for over half of the population. As a primary source of livelihood for millions, the health and productivity of crops are directly linked to the economic well-being of these communities. Despite its importance, agriculture faces numerous challenges, with crop diseases being one of the most severe threats to food security and farmer income. Crop diseases not only reduce the yield but also compromise the quality of agricultural produce, leading to significant financial losses for farmers and impacting the broader agricultural supply chain.

Early and accurate identification of crop diseases is crucial to preventing widespread damage. Effective management strategies depend on timely diagnosis, allowing farmers to take corrective actions before infections spread uncontrollably. Traditionally, crop disease diagnosis involves visual inspection by experts who examine symptoms such as discoloration, spots, lesions, or deformities on leaves, stems, or fruits. However, this manual approach has several limitations. It is often time-consuming, subjective, and dependent on the availability of skilled agronomists, which is scarce in many rural and remote areas. Additionally, many crop diseases present similar symptoms, making it difficult even for experts to distinguish among them accurately. This can lead to misdiagnosis, resulting in the improper use of pesticides or other treatments that may harm the environment or exacerbate the problem.

In recent years, advances in artificial intelligence (AI) and computer vision have revolutionized the field of plant disease diagnosis. By automating the detection and classification of diseases through image analysis, these technologies offer scalable and accessible alternatives to traditional methods. Earlier efforts in this direction relied on conventional image processing techniques that required manual feature extraction, such as color, texture, or shape descriptors. While useful, these methods often struggled with complex real-world images due to variations in lighting, background, and plant species, which limited their accuracy and practical application.

The emergence of deep learning, particularly Convolutional Neural Networks (CNNs), has dramatically improved the effectiveness of automated crop disease diagnosis. CNNs can automatically learn hierarchical features directly from raw image data without explicit manual intervention. This ability allows deep learning models to generalize better across diverse

datasets and complex image variations. However, training such models usually requires large, labeled datasets and significant computational resources, which can be challenging to obtain in agricultural contexts.

To overcome these challenges, transfer learning techniques have become widely adopted. Transfer learning leverages CNN architectures pre-trained on massive datasets like ImageNet, which contains millions of images spanning thousands of classes. By fine-tuning these pre-trained models on smaller, domain-specific agricultural datasets, it is possible to achieve high classification accuracy with fewer training samples. This approach significantly reduces training time and improves the model's robustness to variations in disease appearance and environmental conditions.

In this project, the focus is on using MobileNetV2, a lightweight CNN architecture designed for resource-constrained environments such as mobile and embedded devices. MobileNetV2 is particularly suitable for field applications where processing power and memory are limited. Its design incorporates depthwise separable convolutions and inverted residual blocks, which reduce model size and computational complexity without sacrificing accuracy. These features make MobileNetV2 an excellent candidate for real-time crop disease detection directly on smartphones or low-cost devices.

The model was trained on a comprehensive dataset containing 68,187 images spanning 42 classes, including multiple crops and disease types, as well as healthy leaf conditions. The dataset is diverse, with images captured under varying environmental conditions and backgrounds, which helps the model generalize well to real-world scenarios. To ensure robust evaluation, the dataset was split into training, validation, and testing subsets with balanced class distributions.

Data augmentation techniques such as random rotation, flipping, and scaling were applied during training to improve the model's ability to generalize to new, unseen images. After extensive experimentation and hyperparameter tuning, the MobileNetV2 model achieved an overall accuracy of 92.45% on the test set. In addition to accuracy, metrics such as precision, recall, and F1-score were calculated to assess the model's performance comprehensively, confirming its reliability in distinguishing between healthy and diseased leaves while minimizing false positives and false negatives.

The practical importance of this lightweight, high-accuracy model lies in its integration into a web-based application tailored for use by farmers and agricultural workers. The platform

allows users to upload images of crop leaves taken directly in the field using smartphones or other digital devices. The application quickly processes the uploaded images and provides immediate diagnostic feedback, identifying the presence and type of disease affecting the crop. This real-time detection system empowers users to take prompt action, helping prevent disease spread and reduce crop losses.

Early disease detection enabled by this system is essential for minimizing damage to crops and maintaining yield quantity and quality. Moreover, timely intervention helps avoid the indiscriminate use of pesticides and fungicides, which can have adverse environmental effects such as soil and water contamination, harm to beneficial insects, and the development of pesticide-resistant pathogens. By promoting targeted and responsible agrochemical use, this AI-powered tool supports sustainable agriculture and environmental conservation.

This project also aligns with the broader trend of precision agriculture, which uses data-driven approaches and advanced technologies to optimize resource usage, improve productivity, and reduce environmental impacts. Incorporating automated crop disease diagnosis into this paradigm enables smarter farming decisions and more efficient management of inputs like water, fertilizers, and pesticides. Such technologies are especially valuable in regions with limited access to expert agricultural advice, thereby democratizing agricultural knowledge and empowering smallholder farmers to protect their crops and livelihoods more effectively.

In conclusion, this project presents a scalable and robust deep learning framework based on MobileNetV2 for multi-class classification of crop leaf diseases. The model's strong performance demonstrates its ability to accurately differentiate among various diseases and healthy conditions across diverse real-world scenarios. By harnessing artificial intelligence, this approach addresses critical challenges in agriculture and provides farmers with intelligent tools to enhance crop health management, increase productivity, and improve economic outcomes.

Looking ahead, integrating this disease detection system with IoT, drones, and environmental sensors can enable continuous monitoring and real-time farm decision-making. This will support predictive analytics and proactive disease management, advancing precision agriculture globally. Such innovations will make agriculture more resilient, sustainable, and efficient to meet growing food demands.

CHAPTER 2

PROBLEM STATEMENT AND OBJECTIVES

2.1 Problem Statement

Agriculture plays a vital role in sustaining economies worldwide, especially in developing countries like India, where a large portion of the population depends on farming for their livelihood. However, crop productivity and quality are often threatened by various diseases that cause significant yield losses and economic damage. These diseases reduce both the quantity and quality of agricultural produce, negatively affecting farmers' income and food security.

Traditional disease detection relies heavily on manual visual inspection by agricultural experts who identify symptoms such as spots, discoloration, and deformities on leaves, stems, or fruits. While effective when performed by specialists, this method is time-consuming, labor-intensive, and often unavailable to many farmers, especially in remote or resource-limited areas. Additionally, many diseases exhibit similar symptoms, making accurate diagnosis difficult without expert knowledge. Misdiagnosis can lead to inappropriate treatment and overuse of pesticides, which increases costs and harms the environment.

To overcome these challenges, there is a need for automated, rapid, and accurate disease detection methods that assist farmers in early identification and appropriate treatment. Advances in computer vision and deep learning provide promising solutions by enabling automated analysis of plant images to detect disease symptoms. However, these approaches often require large, labeled datasets, face challenges due to variations in image quality and backgrounds, and must be efficient enough to run on mobile or edge devices commonly used in agricultural fields.

This project addresses the problem of developing a reliable, accurate, and computationally efficient crop disease classification system to assist farmers and agricultural workers in timely diagnosis using leaf images. The system aims to minimize dependence on expert intervention, reduce the time and cost of diagnosis, and adapt to different crops and disease types. It focuses on using a lightweight deep learning model that balances high accuracy with efficiency, suitable for deployment on resource-constrained devices. By improving disease management, this solution supports sustainable agricultural practices and contributes to enhanced food security and farmer livelihoods.

2.2 Objectives

The primary objective of this project is to develop an efficient and accurate system for the detection and classification of crop diseases using deep learning techniques. This system aims to assist farmers, agricultural workers, and researchers in identifying crop diseases at an early stage to minimize crop loss, reduce economic damage, and promote sustainable agricultural practices. To achieve this, the project is guided by the following specific objectives:

1. **Develop a robust dataset:** To compile and utilize a comprehensive dataset consisting of images representing various crop diseases across multiple crop species. The dataset should include images of healthy and diseased leaves, capturing diverse environmental conditions and disease symptoms to enhance the model's generalization capability.
2. **Implement effective image preprocessing:** To design and apply image preprocessing techniques such as resizing, normalization, and augmentation. These steps aim to improve the quality of input images, increase dataset diversity, and enable the model to learn more robust features that improve classification accuracy.
3. **Select an appropriate deep learning model:** To choose a lightweight and efficient convolutional neural network (CNN) architecture suitable for crop disease classification. The model should balance high accuracy with computational efficiency, enabling deployment on resource-constrained devices such as smartphones or edge devices used in agricultural fields.
4. **Train the model using transfer learning:** To leverage transfer learning techniques by fine-tuning a pre-trained CNN model on the prepared crop disease dataset. This approach reduces training time and enhances performance, particularly when the available dataset size is limited compared to large-scale datasets like ImageNet.
5. **Evaluate model performance:** To rigorously evaluate the trained model using relevant metrics such as accuracy, precision, recall, F1-score, and confusion matrix. This evaluation will ensure the model's effectiveness in correctly identifying different crop diseases and healthy conditions.

6. **Develop a user-friendly application:** To integrate the trained model into an accessible platform (e.g., a web or mobile application) that allows users to upload leaf images and receive real-time diagnostic feedback. This tool aims to provide farmers and agricultural workers with an easy-to-use solution for timely disease detection.
7. **Promote sustainable agriculture:** By enabling early and accurate disease identification, the system seeks to reduce the misuse of pesticides and chemical treatments, thus minimizing environmental impact and promoting eco-friendly farming practices.
8. **Identify limitations and future improvements:** To acknowledge the system's current limitations, such as potential challenges in detecting diseases under varying lighting conditions or with limited image quality, and to suggest directions for future enhancements, including expanding the dataset, incorporating additional crop species, and improving model robustness.

CHAPTER 3

LITERATURE SURVEY AND EXISTING SYSTEM

This chapter presents a comprehensive literature survey on crop disease detection using image processing, machine learning, and deep learning techniques. It also describes existing systems, their workflows, and limitations, highlighting the need for advanced solutions.

3.1 Literature Survey

3.1.1 Traditional Image Processing Techniques

Early research relied on classical image processing methods like color thresholding, segmentation, edge detection, and texture analysis.

- *Patil and Kumar (2011)* applied k-means clustering to segment diseased areas in leaves, using RGB histograms and texture features.
- These methods were sensitive to lighting and background changes and required manual feature engineering, limiting their robustness.

3.1.2 Machine Learning-Based Approaches

Machine learning models improved classification by learning from extracted features such as color, shape, and texture.

- *Phadikar and Sil (2008)* used Support Vector Machines (SVM) to detect rice diseases from leaf images with promising accuracy.
- Other models like Random Forests, Naive Bayes, and k-NN were also explored but heavily depended on feature quality and suffered on unseen data.

3.1.3 Deep Learning Revolution

Deep learning, especially Convolutional Neural Networks (CNNs), automated feature extraction and improved accuracy dramatically.

- *Mohanty et al. (2016)* used AlexNet and GoogleNet on PlantVillage dataset achieving over 99% accuracy.
- *Sladojevic et al. (2016)* developed a CNN to classify 13 plant diseases with 96.3% accuracy.
- *Too et al. (2019)* compared ResNet, DenseNet, Inception, and VGG; DenseNet gave the best accuracy but required high computational resources.

3.1.4 Lightweight CNN Models

To enable mobile and edge deployment, lightweight CNNs like MobileNet, ShuffleNet, and EfficientNet-Lite were designed.

- *Howard et al. (2017)* introduced MobileNetV2, which uses depthwise separable convolutions for efficiency without much accuracy loss.
- These models are suitable for offline diagnosis in resource-limited areas but trade-off some accuracy.

3.1.5 Transfer Learning

Transfer learning leverages pre-trained CNNs on large datasets to adapt models for crop disease detection.

- *Ferentinos (2018)* achieved over 99% accuracy on 58 crop-disease classes using transfer learning with VGG16 and ResNet50 backbones.
- Transfer learning reduces training time and data needs, enabling fast adaptation across crops and diseases.

3.1.6 Practical Tools and Applications

Mobile apps like Plantix, Leaf Doctor, and AgroAI use deep learning for real-time disease diagnosis with user-friendly interfaces.

- These apps integrate cloud-based inference and image preprocessing to assist farmers in the field.
- Some platforms also recommend treatment and track disease spread geographically.

3.1.7 Challenges in Existing Systems

- Generalization issues due to diverse field conditions (lighting, occlusions).
- Small, unbalanced datasets limiting model robustness.
- Lack of explainability in deep learning models.
- Trade-offs between speed and accuracy, especially on mobile devices.
- Deployment challenges in rural areas with limited connectivity and hardware.

3.1.8 Other Notable Studies

- *Picon et al. (2019)* used UAV imagery combined with CNNs for large-scale crop health monitoring.
- *Barbedo (2018)* reviewed the use of image processing for plant disease detection, highlighting challenges in real-world applications.
- *Rumpf et al. (2010)* combined hyperspectral imaging with machine learning for early disease detection in cereals.

- *Mohanty and Hughes (2018)* explored data augmentation to improve deep learning model robustness.
- *Singh et al. (2020)* proposed hybrid CNN-SVM models to enhance classification accuracy.

3.2 Existing System

3.2.1 Traditional Manual Diagnosis

Farmers visually inspect plants for disease symptoms, which is slow, subjective, and error-prone, especially for large farms or non-expert farmers.

3.2.2 Automated Disease Detection Systems

Modern systems automate diagnosis through a pipeline:

1. Image Acquisition
 - Images are captured by smartphones or cameras, where image quality is vital.
2. Image Preprocessing
 - Techniques like resizing, noise filtering, and color space conversion (RGB to HSV) standardize images.
3. Feature Extraction
 - Classical methods extract color, texture, and shape features to represent disease symptoms.
4. Classification
 - Extracted features are input into classifiers like SVM, k-NN, or Random Forest for disease categorization.




Figure 3.2: Workflow of Existing Plant Disease Detection System

3.2.3 Limitations of Existing Systems

- Dependence on manual feature engineering.
- Sensitivity to environmental factors.
- Poor generalization across crops and conditions.
- Computationally expensive and slow.
- Lack of support for real-time or mobile use.

3.2.4 Need for Advanced Solutions

Traditional image processing and machine learning techniques, while useful, often require manual feature extraction and are highly sensitive to variations in lighting, background, and leaf orientation. These limitations hinder their reliability and scalability in real-world agricultural scenarios. To address these challenges, deep learning models—particularly Convolutional Neural Networks (CNNs)—have emerged as a promising alternative due to their ability to automatically learn hierarchical features directly from image data. These models demonstrate significantly higher accuracy and robustness, making them ideal for complex classification tasks such as identifying multiple crop diseases across diverse species.

Furthermore, recent developments in lightweight CNN architectures like MobileNet, EfficientNet, and SqueezeNet have enabled deployment on mobile and edge devices. This facilitates real-time disease detection even in remote farming areas with limited computational resources. These mobile-compatible models ensure scalability, reduce the need for internet connectivity, and empower farmers with on-field diagnostic tools. Therefore, there is a critical need to shift towards these advanced deep learning solutions to enhance precision agriculture, reduce crop loss, and support food security through early and accurate disease diagnosis.

CHAPTER 4

PROPOSED SYSTEM AND ARCHITECTURE

4.1 Proposed System Overview

The proposed system aims to provide a real-time, scalable, and accurate solution for crop leaf disease classification using deep learning techniques. Specifically, it leverages the MobileNetV2 convolutional neural network, which is renowned for its balance of speed and accuracy, particularly on mobile and low-resource devices.

This system addresses the shortcomings of traditional image processing and machine learning techniques, which rely on handcrafted features and are sensitive to environmental variations such as lighting conditions, leaf positioning, and background clutter. Instead, MobileNetV2 learns directly from raw image pixels, offering more robustness and automation in feature extraction and classification.

Through an easy-to-use web or mobile interface, farmers can capture and upload images of diseased crop leaves. The system processes the image, classifies the disease into one of 42 categories (including healthy leaves), and provides immediate feedback with suggested actions. This functionality empowers farmers with real-time insights for timely disease management.

4.2 Key Components of the Proposed System

1. Image Upload Module

- Allows users to upload or capture leaf images using smartphones or a web interface.
- Supports uploads from local storage for flexibility in usage.

2. Preprocessing Module

- Resizes input images to 224×224 pixels to match MobileNetV2's input requirements.
- Normalizes pixel values for consistent model behavior.

3. Model Inference Engine (MobileNetV2)

- Loads a pre-trained MobileNetV2 model.

- Performs classification to predict one of 42 disease classes, including healthy leaves.

4. Result Interpretation Module

- Converts prediction into a readable label (e.g., “Tomato - Early Blight”).
- Displays a short description and recommended actions.
- Shows probability/confidence scores to support decision-making.

5. Feedback and Logging Module

- Allows users to provide optional feedback on predictions.
- Stores logs for performance analysis and future model updates.




Fig4.1:Proposed system architecture

4.3 Benefits of the Proposed System

- Real-Time Diagnosis: Enables immediate detection to prevent disease spread.
- Mobile-Friendly: Optimized for mobile devices, facilitating field-level use.
- Scalable: Easily adaptable to more crops and diseases.
- High Accuracy: Achieves ~92.45% accuracy across 42 disease and healthy classes.
- User-Friendly: Designed for ease of use by farmers and agricultural workers.

4.4 Why MobileNetV2?

MobileNetV2 is a lightweight convolutional neural network designed for efficient performance on mobile and embedded systems. It uses depthwise separable convolutions and inverted residual blocks, which drastically reduce the number of parameters and computational cost while retaining classification accuracy. These attributes make it ideal for use in rural and remote settings, where computational resources and internet connectivity may be limited. MobileNetV2's efficiency allows it to be deployed via smartphones, ensuring accessibility for smallholder farmers.

4.5 Deployment Strategy

- **Model Training:** Conducted on GPU-enabled platforms with a dataset of 68,187 images spanning 42 classes.
- **Deployment Platform:** Utilizes Flask for web deployment, with mobile compatibility via Streamlit with PWA or React Native for app integration.
- **User Interface:** Simple interface enabling image upload, prediction display, and feedback collection.

4.6 System Architecture

The system architecture integrates all functional components, streamlining the process from image acquisition to prediction output and feedback.

1. Input Layer – Leaf Image Collection

- Images are captured via camera or uploaded from device storage.
- Organized into folders based on disease categories for model training and evaluation.

2. Preprocessing Module

- Resizes images to 224×224 pixels.
- Normalizes pixel values to a [0,1] range.
- Applies data augmentation (rotation, zoom, shift, flip) to improve model generalization.

3. Feature Extraction – MobileNetV2

- A pre-trained MobileNetV2 (with top layer removed) extracts deep features from input images.
- Reduces computational cost while capturing detailed disease-related patterns.

4. Classification Head

- Custom dense layers are appended:
- `Dense(1024, activation='relu') → Dense(512, activation='relu') → Dense(NUM_CLASSES, activation='softmax')`
- Outputs probability scores for each class.

5. Evaluation Module

- Model is evaluated on both training and testing datasets.
- Metrics include accuracy, loss, and confusion matrix for performance validation.

6. Deployment Layer

- The final trained model is saved in .h5 format.
- Deployed via Streamlit web app for real-time inference.
- Includes:
 - Image upload interface.
 - Immediate prediction and feedback.
 - Logging for continuous improvement.




Fig4.6: System architecture

CHAPTER 5

DATASET DESCRIPTION AND IMAGE PREPROCESSING

5.1 Dataset Description

5.1.1 Overview of Dataset

- Crop leaf images across 42 categories (diseased and healthy)
- Diverse environmental conditions and visual variations

5.1.2 Dataset Composition

- Total Images: 69,187
- Training Set: 28,830
- Testing Set: 28,830
- Validation Set: 11,527

5.1.3 Image Resolution

- All images resized to **224x224 pixels** for deep learning model compatibility

5.1.4 Data Augmentation Summary

- Techniques: Rotation, Flipping, Zooming
- Purpose: To improve generalization in real-world conditions.




Figure 5.1: Sample Images from the dataset

5.2 Image Preprocessing Techniques

5.2.1 Importance of Preprocessing

- Mitigates inconsistencies in field-captured images
- Focuses model on disease-specific patterns, not background noise

5.3 Preprocessing Steps

5.3.1 Image Resizing and Scaling

- Resized to 224x224 pixels
- Pixel normalization from $[0, 255]$ to $[0, 1]$

5.3.2 Image Denoising and Enhancement

- **Median Filtering:** Removes salt-and-pepper noise
- **Gaussian Blur:** Smoothens image and removes sharp edges
- **Histogram Equalization:** Enhances contrast
- **Edge Detection/Thresholding:** Highlights affected regions

5.3.3 Data Augmentation

- **Rotation** ($\pm 15^\circ$ to $\pm 45^\circ$)
- **Flipping** (horizontal and vertical)
- **Zooming and Shifting**
- **Brightness Adjustment**
- **Shearing** to mimic natural leaf curvature

Implemented using TensorFlow's **ImageDataGenerator**

5.4 Color Space Transformations

5.4.1 RGB as Primary Color Space

- Standard format used during training

5.4.2 HSV and LAB for Feature Analysis

- HSV: Better color change visualization (e.g., fungus)
- LAB: Contrast and brightness separation

5.5 Image Format and Label Management

5.5.1 Format Conversion

- All images stored in JPEG or PNG
- Removed blank or unreadable images

5.5.2 Label Encoding

- Converted text labels to numeric form (e.g., apple_scab → 0)
- Essential for model compatibility and output mapping

5.6 Dataset Directory Structure

5.6.1 Hierarchical Organization

The dataset is organized in a folder-based hierarchy to simplify training and labeling. Deep learning libraries like Keras and PyTorch use folder names to auto-assign labels.

- Training Set: Images are stored inside folders like
dataset > train > class_1,
dataset > train > class_2, etc.
- Validation Set: Images are stored in folders like
dataset > validation > class_1,
dataset > validation > class_2, etc.
- Test Set: Images are stored in folders such as
dataset > test > class_1,
dataset > test > class_2, etc.

Each *class_x* folder contains images corresponding to that specific class.

5.6.2 Framework Compatibility

- Enables automatic label loading in Keras/PyTorch

5.6.3 Balanced Class Distribution

- Each subset includes all 42 classes

5.7 Handling Dataset Imbalance

5.7.1 Class Distribution Issues

- Some classes had fewer samples

5.7.2 Targeted Augmentation

- More transformations applied to minority classes

5.7.3 Class Weights in Loss Function

- Heavier penalties for misclassifying underrepresented classes

5.7.4 Evaluation Metrics

- Used Precision, Recall, and F1-Score instead of accuracy

5.8 Significance of Preprocessing

- Promotes uniformity
- Removes noise
- Highlights disease-specific patterns
- Improves model generalization
- Enhances accuracy and fairness across all classes

CHAPTER 6

MODEL SELECTION AND ARCHITECTURE

6.1 Model Selection – MobileNetV2

Selecting the appropriate model architecture is critical for building an efficient and accurate deep learning system for crop disease classification. Among various options, MobileNetV2 stands out due to its excellent trade-off between accuracy and computational efficiency. Developed by Google in 2018, this architecture is tailored for mobile and embedded devices, making it highly suitable for agricultural applications in resource-constrained environments such as smartphones or drones.

6.2 Overview of MobileNetV2 Architecture

MobileNetV2 improves upon MobileNetV1 by introducing two key innovations:

- **Inverted Residual Blocks:**
 - Expand input channels using 1x1 convolution.
 - Apply depthwise separable convolutions for spatial processing.
 - Project back to lower dimensions, reducing computation while preserving information.
- **Linear Bottlenecks:**
 - Use linear activations in bottleneck layers to minimize information loss.
 - Essential during dimensionality reduction to preserve feature integrity.

The overall architecture includes:

- Initial convolution layer
- Multiple inverted residual blocks with varying expansion and stride
- Final convolution, global average pooling, and a fully connected layer

6.3 Why MobileNetV2 is Suitable for Crop Disease Classification

1. Efficient and Lightweight

MobileNetV2 is compact (<15MB) and optimized for real-time inference on devices with limited power and memory.

2. Effective Transfer Learning

Using pre-trained ImageNet weights accelerates training and improves generalization on limited crop disease datasets.

3. Robust to Field Variations

Handles variability in lighting, background, and orientation using depthwise convolutions and bottleneck designs.

4. Flexible Input Size

Supports various input resolutions (default: 224x224) with minimal performance trade-offs.

6.4 Technical Specifications of MobileNetV2

Feature	Value
Parameters	~3.4 million
MACs (Multiply-Adds)	~300 million (for 224x224 input)
Activation Functions	ReLU6 (intermediate), Linear (bottlenecks)
Convolutions	Depthwise Separable
Pre-trained Dataset	ImageNet

6.5 Training Strategy for Crop Disease Classification

Transfer Learning: Initialize with ImageNet weights

Freeze Layers: Freeze early layers to preserve general features

Fine-Tune: Train deeper layers and classification head

Data Augmentation: Apply rotations, zooms, flips, and lighting variations

Optimizers: Use Adam or RMSprop with learning rate scheduling

6.6 Comparison with Other Models

Model	Parameters	Accuracy	Suitability for Edge
VGG16	~138M	High	Too heavy
ResNet50	~25M	Very High	Moderate
MobileNetV2	~3.4M	High	Excellent

Compared to heavier models like ResNet50 and VGG16, MobileNetV2 offers competitive accuracy with significantly reduced size and computational requirements. While some deeper networks may achieve slightly higher accuracy, MobileNetV2's speed and efficiency make it more practical for on-device deployment in agricultural environments.

6.7 Limitations of MobileNetV2

- May not match the performance of very deep or ensemble models on highly complex datasets.
- Still requires careful **data preprocessing, augmentation, and hyperparameter tuning** for optimal results.

6.8 Model Architecture

MobileNetV2 is an efficient and accurate model designed for image classification on resource-limited devices. It processes an input image of size **224x224x3** through convolutional and residual blocks to extract features.

Input Layer:

The input image size is **224x224x3**.

Initial Convolution (Conv1):

Reduces spatial size to **112x112** and increases channels to **32** for early feature extraction.

Residual Blocks:

Uses inverted residual blocks with depthwise separable convolutions and shortcuts:

- **Block 1:** Outputs 112x112x16, preserving resolution and reducing channels.
- **Block 2:** Maintains 112x112, channels increase to 24.
- **Block 3:** Downsamples to 56x56, channels to 32.
- **Block 4:** Downsamples to 28x28, channels to 64.
- **Block 5:** Downsamples to 14x14, channels to 96.
- **Block 6:** Downsamples to 7x7, channels to 160, capturing high-level features.

Final Convolution:

Generates 7x7x1280 feature maps to enrich feature representation.

Global Average Pooling:

Reduces to 1x1x1280, summarizing global features and reducing parameters.

Output Layer (Softmax):

Fully connected layer predicts probabilities over 1000 classes (adjustable).

6.9 Key Highlights:

- **Inverted Residuals:** Expansion, depthwise convolution, and projection optimize extraction.
- **Depthwise Separable Convolutions:** Lower computational cost with maintained accuracy.
- **Progressive Downsampling:** Extracts local and global features efficiently.
- **Global Average Pooling:** Reduces overfitting and parameters.

This design balances efficiency and accuracy, ideal for mobile and edge applications such as crop disease detection

.




Figure 6.1: MobileNetV2 Architecture

CHAPTER 7

MODEL TRAINING, METRICS, RESULTS, AND COMPARISION

7.1 Model Training

Model training is the cornerstone of any deep learning pipeline. For this study, we trained the MobileNetV2 model using transfer learning and fine-tuning strategies. The process was carefully crafted to adapt the model to the specific characteristics of crop disease datasets, which often include diverse background noise, lighting variations, and intra-class variability.

Data Preprocessing and Augmentation

Before training, all images were resized to 224x224 pixels to match the input dimension of MobileNetV2. The dataset was then divided as follows:

- 50% for training
- Separate dataset for validation
- 50% for testing

To prevent overfitting and improve generalization, data augmentation techniques were applied:

- Random rotations (± 30 degrees)
- Horizontal and vertical flipping
- Brightness and contrast adjustments
- Zooming (0.8x–1.2x range)
- Shearing and translation

Transfer Learning Setup

MobileNetV2 was initialized with ImageNet-pretrained weights to leverage generalized visual features such as edges, patterns, and textures. Initially, all layers were frozen except the classification head.

Fine-tuning Strategy

After initial training, the last few convolutional blocks were unfrozen and fine-tuned using a lower learning rate. This allowed the model to adjust deeper representations based on crop-specific patterns like leaf discoloration, disease spots, and fungal textures.

Training Configuration

- Loss Function: Categorical Crossentropy
- Optimizer: Adam
- Batch Size: 32
- Epochs: 50 (with early stopping based on validation loss)
- Learning Rate Scheduler: ReduceLROnPlateau
- Callbacks: EarlyStopping, ModelCheckpoint

This training strategy provided a balanced mix of speed, generalization, and performance, enabling robust classification even on challenging crop disease images.

7.2 Evaluation Metrics

To evaluate model performance, multiple metrics were used. Each offers a different perspective on the model's behavior.

Accuracy

Accuracy is the proportion of correctly classified instances:

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

It gives a general sense of model performance but may be misleading if the dataset is imbalanced.

Precision, Recall, and F1-Score

- **Precision** evaluates how many predicted positives are truly positive.
- **Recall** assesses how many actual positives were identified.
- **F1-Score** is the harmonic mean of precision and recall, especially useful for class imbalance.

$$\text{F1 Score} = 2 \times \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Confusion Matrix

A confusion matrix was generated to visualize class-wise performance. It shows true labels vs. predicted labels, helping identify specific misclassification trends (e.g., diseases with similar visual features).

ROC Curve and AUC

Receiver Operating Characteristic (ROC) curves were plotted for multi-class classification. The Area Under the Curve (AUC) summarizes the trade-off between the true positive rate and false positive rate across thresholds.

Top K Accuracy

Top 3 and Top 5 accuracy were also reported to evaluate how often the correct class appeared in the top predictions, which is useful in real-world applications where multiple possibilities can be presented to the user.

7.3 Result Analysis

Training and Validation Performance

MobileNetV2 achieved:

- Training Accuracy: ~93.6%
- Validation Accuracy: ~92.4%
- Test Accuracy: 91.8%

The training and validation loss curves indicate smooth convergence without significant overfitting. Early stopping was triggered after ~37 epochs, suggesting an optimal balance between underfitting and overfitting.

Error Analysis

- Poor lighting conditions in some images affected performance.
- Blurry or occluded leaves reduced feature clarity.
- These issues can be improved with better image collection protocols or by incorporating denoising techniques.

7.4 Comparison with Existing Models

To assess MobileNetV2's practical value, its performance was benchmarked against three popular deep learning models: ResNet50, VGG16, and EfficientNetB0.

Model	Accuracy	Strengths	Weaknesses
VGG16	~93%	Deep feature extraction, strong generalization	Very large model size (>500MB), high inference latency
ResNet50	~94%	Residual connections ensure robust performance	Computationally heavy, slower inference on edge/mobile devices
DenseNet121	~95%	Dense connectivity improves feature reuse and accuracy	Demands powerful hardware and large memory
MobileNetV2	92.45%	Lightweight, fast, ideal for mobile and embedded devices	Slightly reduced accuracy for visually similar disease types
SVM + HOG	~75–80%	Simple, interpretable, low-resource usage	Poor generalization, manual feature extraction required

Performance Insights

- MobileNetV2 delivered competitive accuracy while being extremely lightweight.
- ResNet50 slightly outperformed in accuracy but is heavier and slower, not ideal for field deployment.
- VGG16 had the highest parameter count, leading to memory and speed issues.
- EfficientNetB0 offered a good balance, but MobileNetV2 still excelled in speed and deployability.

CHAPTER 8

System Advantages, Limitations, and Future Scope

8.1 Advantages of the Proposed System

The proposed MobileNetV2-based crop disease classification system demonstrates a range of features that collectively offer immense utility for the agricultural sector, particularly in crop protection and disease management.

8.1.1 High Accuracy

The model delivers an outstanding **accuracy of 92.45%** across **42 different crop disease classes**, which indicates a strong ability to distinguish between various types of infections and healthy crops. This high performance facilitates **early detection**, allowing timely remedial measures and minimizing crop damage. Such precision greatly improves farmers' confidence in technology-based diagnosis over manual observation.

8.1.2 Mobile-Friendly Deployment

MobileNetV2 is specially designed for low-latency and low-computation environments. Its lightweight architecture makes it ideal for deployment on mobile phones, tablets, and edge devices, eliminating the need for powerful GPUs or cloud infrastructure. As a result, farmers in rural or under-resourced areas can also benefit from advanced AI tools without expensive hardware.

8.1.3 Real-Time Inference

The system's real-time response feature means users get near-instant feedback upon uploading an image. Such immediate results are vital in agricultural contexts where rapid decision-making can mean the difference between crop salvage and loss.

8.1.4 User-Friendly Interface

A simple, intuitive web-based interface makes the system accessible to non-technical users. Farmers can easily take pictures of leaves using their mobile devices and upload them to the system. The results are presented in a clear, actionable format without requiring deep understanding of machine learning.

8.1.5 Support for Sustainable Agriculture

This system aligns with sustainable farming practices. By facilitating accurate and timely diagnosis, it helps reduce the overuse of pesticides and fertilizers, thereby minimizing environmental harm and cutting unnecessary costs. The system also promotes targeted intervention, leading to better yield outcomes with fewer resources.

8.1.6 Scalability and Extensibility

Thanks to its modular design, the model can be easily fine-tuned to support **new crop** species or emerging diseases. As more data becomes available, the system can be updated without reengineering the entire architecture, allowing it to evolve with agricultural needs.

8.1.7 Robustness to Variability

Using advanced image preprocessing and data augmentation techniques like rotation, flipping, brightness changes, and zooming, the model is trained to handle a wide range of input variations. This ensures reliable performance across diverse environmental conditions, image qualities, and camera devices.

8.1.8 Efficient Training via Transfer Learning

By utilizing transfer learning with a pre-trained ImageNet MobileNetV2 model, the system reduces the need for large datasets and extensive training time. This approach enables rapid prototyping and deployment with limited resources while maintaining high accuracy.

8.1.9 Strong Generalizability

The model was trained on a diverse dataset that closely mimics real-world conditions, allowing it to generalize effectively across unseen samples. This ensures that the system remains useful and reliable not only in lab settings but also in actual field deployments.

8.1.10 Cloud and Edge Compatibility

Depending on infrastructure availability, the system can be deployed on both cloud platforms and on-device edge computing systems. This flexibility ensures continuous operation, even in offline or low-connectivity scenarios common in remote farms.

8.2 Limitations

While the system shows strong performance and potential, several limitations must be acknowledged to guide future improvements and realistic deployment expectations.

8.2.1 Dataset Limitations

Despite training on over 68,000 images, the dataset might not represent all real-world scenarios, such as rare diseases, specific regional plant strains, or symptoms affected by non-disease stressors like nutrient deficiency or pesticide exposure. This can lead to misclassification or lower accuracy in unrepresented conditions.

8.2.2 Image Quality Sensitivity

The model's performance is heavily dependent on image clarity and focus. Images taken in poor lighting, extreme weather, or improper angles can result in incorrect predictions. This is especially challenging when using low-end mobile cameras.

8.2.3 Lack of Multi-Disease Detection

In real-world farming, it is common for plants to suffer from multiple diseases simultaneously. The current model tends to identify only the dominant disease, which may lead to incomplete diagnoses and ineffective treatment plans.

8.2.4 Black-Box Nature

As with most deep learning models, the system operates as a black box, offering minimal insight into the decision-making process. This lack of interpretability may hinder user trust and make it difficult to validate or cross-check predictions.

8.2.5 Device and Connectivity Limitations

Although optimized for mobile deployment, low-end smartphones may still face performance lags. Additionally, if internet connectivity is required, areas with limited or no internet access may not benefit from the system unless offline modes are developed.

8.2.6 Usability Barriers

The current web application may be less accessible to users with low digital literacy or those speaking regional languages. Without features like voice support, local dialect translation, or interactive tutorials, the system may see limited adoption among its target users.

8.2.7 Risk of False Predictions

Despite high overall accuracy, the risk of false positives or false negatives remains. This can result in unnecessary pesticide application or missed disease detection, both of which negatively impact crop health and farmer trust.

8.2.8 Inconsistent Image Input

There are no standardized protocols for how users should capture images. This can lead to inconsistent input quality, affecting the model's reliability. For example, leaf orientation, lighting, and background can introduce unwanted variability.

8.2.9 Limited Feedback Mechanism

The absence of an integrated user feedback loop hinders the model's ability to learn from real-world errors. Continuous improvement depends on active feedback from users, which is currently not utilized.

8.2.10 Data Privacy Concerns

As the system involves capturing and transmitting agricultural data, including farm locations and conditions, data privacy and security become critical concerns that must be addressed through encryption and regulatory compliance.

8.3 Future Scope

Looking ahead, the system offers multiple pathways for innovation and enhancement, making it a foundational step toward smart and sustainable agriculture.

8.3.1 IoT Integration

Future iterations could include Internet of Things (IoT) sensors and cameras across fields to monitor crop health in real time. These sensors can automatically capture images and environmental data, triggering early alerts when disease symptoms are detected.

8.3.2 Drone-Based Surveillance

The use of drones with HD cameras can scale the system's applicability to large farms,

capturing aerial views for disease mapping. This facilitates region-wide monitoring with minimal human intervention.

8.3.3 Multi-Disease and Severity Detection

Advanced models could be trained to identify multiple infections in a single image and determine the severity level of each disease. This enables more precise treatment plans, reducing chemical use and boosting yield.

8.3.4 Offline Capabilities

To ensure broader adoption, especially in rural or remote areas, the system should support offline deployment on mobile devices. This involves embedding the entire model and interface into an Android/iOS application that runs without internet access.

8.3.5 Regional Language and Voice Interface

Integrating regional language support and voice-based navigation would make the system accessible to non-literate and native-speaking farmers, eliminating language barriers and enhancing usability.

8.3.6 Explainable AI (XAI)

Incorporating XAI techniques like Grad-CAM or LIME would help visualize the regions of the image that influenced the model's decision, offering transparency and increased user trust.

8.3.7 Mobile App Integration

A full-featured mobile app could combine disease detection, treatment suggestions, weather alerts, market price info, and government schemes—creating a comprehensive digital farming assistant.

8.3.8 Policy and Data Integration

By integrating with government agriculture portals and databases, the system could provide real-time crop health maps, aiding in outbreak control, policy formulation, and resource distribution during epidemics.

8.3.9 Climate and Soil Integration

Long-term, the model could be enhanced with data on climate patterns, soil health, and crop history, enabling predictive features like disease risk forecasting and planting schedule recommendations.

8.3.10 Continuous Improvement through Collaboration

Establishing a collaborative ecosystem involving farmers, agronomists, researchers, and technologists will drive ongoing improvements. Such partnerships will ensure that the system remains updated, relevant, and beneficial to all stakeholders.

CHAPTER 9

REFERENCES

1. Mohanty, S. P., Hughes, D. P., & Salathé, M. (2016). Using deep learning for image-based plant disease detection. *Frontiers in Plant Science*, 7, 1419. <https://doi.org/10.3389/fpls.2016.01419>
2. Howard, A. G., et al. (2017). MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications. *arXiv preprint arXiv:1704.04861*. <https://arxiv.org/abs/1704.04861>
3. Ferentinos, K. P. (2018). Deep learning models for plant disease detection and diagnosis. *Computers and Electronics in Agriculture*, 145, 311–318. <https://doi.org/10.1016/j.compag.2018.01.009>
4. Too, E. C., Yujian, L., Njuki, S., & Yingchun, L. (2019). A comparative study of fine-tuning deep learning models for plant disease identification. *Computers and Electronics in Agriculture*, 161, 272–279. <https://doi.org/10.1016/j.compag.2018.03.032>
5. Brahimi, M., Boukhalfa, K., & Moussaoui, A. (2017). Deep learning for tomato diseases: Classification and symptoms visualization. *Applied Artificial Intelligence*, 31(4), 299–315. <https://doi.org/10.1080/08839514.2017.1315516>
6. TensorFlow. (n.d.). TensorFlow Lite. Retrieved from <https://www.tensorflow.org/lite>
7. Kaggle. (n.d.). PlantVillage Dataset. Retrieved from <https://www.kaggle.com/datasets/emmarex/plantdisease>
8. Simonyan, K., & Zisserman, A. (2014). Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556*. <https://arxiv.org/abs/1409.1556>
9. Deng, J., et al. (2009). ImageNet: A large-scale hierarchical image database. *2009 IEEE Conference on Computer Vision and Pattern Recognition*, 248–255. <https://doi.org/10.1109/CVPR.2009.5206848>
10. Barbedo, J. G. A. (2013). Digital image processing techniques for detecting, quantifying and classifying plant diseases. *SpringerPlus*, 2(1), 660. <https://doi.org/10.1186/2193-1801-2-660>
11. Ganesh B. Singh, Rajneesh Rani(2020). Agricultural Crops Disease Identification and Classification through leaf Images using Machine Learning and Deep Learning technique: A review.
12. Jung et al.,(2023), Construction of deep learning- based disease detection model in plants.

CHAPTER 10

CONCLUSION AND APPENDIX

10.1 Conclusion

The deep learning-based crop disease classification system using MobileNetV2 marks a pivotal advancement at the intersection of agriculture and artificial intelligence. As global food demand rises, and challenges like climate change, unpredictable weather, and plant diseases threaten food security, the need for smart, automated, and accessible solutions becomes not just innovative but essential.

This project addresses that need by leveraging convolutional neural networks (CNNs)—specifically the lightweight and efficient MobileNetV2 architecture—to identify crop diseases from leaf images. Through extensive training on a diverse dataset, our model has achieved strong performance in classifying 38 disease types across 14 plant species, while also differentiating them from healthy crops. Its high accuracy, compact structure, and mobile compatibility make it a practical tool for real-world agricultural use, especially in remote and resource-limited regions.

A notable strength of the system lies in its user-friendly web interface. Farmers can upload a plant leaf image and receive an AI-powered diagnosis instantly. This reduces dependence on agricultural experts and empowers farmers to take timely, informed action. Its optimization for mobile devices ensures accessibility, making it scalable from individual use to large-scale deployment by agricultural organizations.

Beyond its technical implementation, the project enriched our learning in:

- Real-time problem solving
- Data preprocessing
- Model evaluation
- Practical deployment

These experiences deepened our understanding of AI's application in real-world scenarios.

While promising, the system has limitations:

- Cannot detect multiple diseases in one image
- Sensitive to poor-quality images
- Lacks offline capability and regional language support

Rather than shortcomings, we view these as opportunities for future enhancement.

In the long term, integration with GPS, climate sensors, and weather APIs can make the system predictive—warning users about disease outbreaks before they occur. Integration with soil health and crop history data could enable precision farming practices that boost yield and sustainability.

This tool can also serve as an educational resource for agricultural students and extension workers, especially when paired with Explainable AI (XAI) techniques like **Grad-CAM**, which clarify how AI models make decisions. Deployment in automated farm machinery—such as sprayers or robotic harvesters—could further revolutionize disease control and harvesting efficiency.

From a societal perspective, this system democratizes access to agricultural knowledge. It empowers smallholder farmers, encourages data-driven farming, and promotes self-reliance in managing crop health.

Importantly, we ensured that this AI tool augments human expertise rather than replaces it. Farmers remain central to decision-making, fostering trust and long-term adoption.

In conclusion, this MobileNetV2-based crop disease classification system is more than just a technical innovation—it's a steppingstone toward smart and inclusive agriculture. With enhancements like multilingual support, offline use, broader datasets, and mobile app integration, this system has the potential to become a global farming assistant. This journey has transformed us not only into developers but also into socially responsible technologists. As we move forward, we are inspired to continue refining our solution and using AI to make agriculture smarter, more resilient, and more accessible.

10.2 Appendix

This appendix contains supporting visuals and code excerpts from the MobileNetV2-based crop disease classification system.

10.2.1 Code Snippets

Model Architecture



```
NUM_CLASSES = len(train_generator.class_indices)
(variable) base_model: Any
base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=(224, 224, 3))
base_model.trainable = False # Freeze base model

x = base_model.output
x = GlobalAveragePooling2D()(x)
x = Dense(1024, activation='relu')(x)
x = Dense(512, activation='relu')(x)
output_layer = Dense(NUM_CLASSES, activation='softmax')(x)

model = Model(inputs=base_model.input, outputs=output_layer)
```

[7] Python

... Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_9406464/9406464 1s 0us/step

Figure 10.2.1: Model Architecture

Data Preprocessing and Augmentation

```
train_datagen = ImageDataGenerator(  
    rescale=1.0/255.0,  
    rotation_range=30,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    shear_range=0.2,  
    zoom_range=0.2,  
    horizontal_flip=True,  
    fill_mode='nearest'  
)  
  
test_datagen = ImageDataGenerator(rescale=1.0/255.0)
```

[5] Python


Figure 10.2.1: Data Preprocessing and Augmentation

Model Training

```
Training Epoch 1/10  
/usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122  
self._warn_if_super_not_called()  
901/901 ————— 380s 409ms/step - accuracy: 0.6546 - loss: 1.3423  
901/901 ————— 119s 129ms/step - accuracy: 0.9015 - loss: 0.3242  
✓ Epoch 1: Train Acc = 0.7789, Test Acc = 0.8760  
  
Training Epoch 2/10  
901/901 ————— 316s 348ms/step - accuracy: 0.8764 - loss: 0.3749  
901/901 ————— 55s 61ms/step - accuracy: 0.9129 - loss: 0.2610  
✓ Epoch 2: Train Acc = 0.8796, Test Acc = 0.8879  
  
Training Epoch 3/10  
901/901 ————— 307s 338ms/step - accuracy: 0.8998 - loss: 0.3062  
901/901 ————— 56s 62ms/step - accuracy: 0.9058 - loss: 0.2794  
✓ Epoch 3: Train Acc = 0.8986, Test Acc = 0.8863  
  
Training Epoch 4/10  
901/901 ————— 307s 338ms/step - accuracy: 0.9041 - loss: 0.2763  
901/901 ————— 56s 62ms/step - accuracy: 0.9320 - loss: 0.1985
```

Figure 10.2.1: Model Training Code

10.2.2 Confusion Matrix



10.2.3 Graphs

Precision, Recall, and F1-Score per Class




Figure 10.2.3: Class-wise Precision, Recall, F1-Score

10.2.4 Underfitting Graph




Figure 10.2.4: Underfitting Curve

10.2.5 Training and Test Accuracy & Loss Curve

Accuracy Curve




Figure 10.2.5: Training vs. Test Accuracy Curve

Loss Curve




Figure 10.2.5: Training vs. Test Loss Curve

10.2.6 Model Accuracy: Correct vs. Total Images



Figure 10.2.6: Bar Chart – Correctly Classified vs. Total Images

10.2.7 User Interface Screenshots

Home Screen of Web Application




Figure 10.2.7: Web App Home Interface

Image Upload and Result Display




Figure 21.9.2: Image Upload and Prediction Output

